

# IoT: Client Devices

Daemons in C

# OS Services

## KERNEL SERVICES, OS SERVICES

---

- ▶ Daemons have access to Kernel and OS services
- ▶ Networking, system logging, processes, filesystem

## THINGS DAEMONS DON'T HAVE

---

- ▶ A console!
- ▶ A user!
- ▶ A home directory!
- ▶ User interaction!

# Console Management

## DAEMONS DON'T HAVE A CONSOLE

---

- ▶ So what do stderr, stdout, and stdin mean?
- ▶ Nothing - you need to manage

## CLOSE STANDARD FILE DESCRIPTORS

---

```
close(STDIN_FILENO) ;  
close(STDOUT_FILENO) ;  
close(STDERR_FILENO) ;
```

# Signal Management

## CONTROLLING INTERACTIVE PROGRAMS

---

- ▶ Unixes give ctrl-c, ctrl-z, etc.
- ▶ Not daemons!
- ▶ Remember, these just submit *signals to processes* (see: **man kill**)

## SIGNAL MANAGEMENT

---

```
signal(SIGKILL, _signal_handler);  
signal(SIGTERM, _signal_handler);  
signal(SIGHUP, _signal_handler);
```

# Logging

NO CONSOLE -> NO STANDARD OUTPUT

---

- Many programs will log to STDERR, or STDOUT
- But we closed them!

SYSLOG

---

- Syslog is a systems-wide logger
- /var/log/messages or /var/log/syslog

OPEN A LOG AND LOG TO IT

---

- `openlog()`, `syslog()`, `closelog()`
- see: **man syslog**

# Working Directory

NO USER -> NO DEFAULT WORKING DIRECTORY

---

- ▶ We need a working directory
- ▶ We do handle files

MOVING AND SETTING A WORKING DIRECTORY

---

```
chdir (WORKING_DIR) ;
```

# File Creation

## WE DO CREATE FILES

---

- ▶ Need to set default permissions on created files
- ▶ Usually files only read by privileged users
- ▶ In our case, better to leave open

## SETTING DEFAULT PERMISSIONS

---

```
umask (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
```

# Sessions

A *SESSION* HAS ONE OR MORE *PROCESS GROUPS*

---

- ▶ The first process in the session is default session leader

A *PROCESS GROUP* HAS ONE OR MORE *PROCESSES*

---

- ▶ The group leader and child processes

THINK OF SESSIONS AS *TERMINAL SESSIONS*

---

**setsid()**



# Forking

## DON'T LOCK UP SPAWNING PROCESS

---

- Daemons need their own dedicated processes
- forking a process creates a copy of the program in another process
- Parent process gets PID; child gets 0; error is negative

## USING FORK()

---

```
PID_T  PID  =  FORK ( ) ;  
IF  (PID > 0)  EXIT (0) ;  
IF  (PID < 0)  EXIT (1) ;
```

# IoT: Client Devices

A Sample Daemon

```
51
52 int main(void) {
53     // Open syslog. We want to write directly to the log with
54     // no wait, and we want to include the PID of the daemon process.
55     // We are running a daemon, so set the LOG_DAEMON flag too.
56     //
57     // Why open here? well, if we have an error forking, we want to
58     // register the error to syslog. To do that, we need an open log.
59     // The child process will inherit this log, as we don't close it
60     // in the parent when it exits, this opening here is A+.
61     openlog(DAEMON_NAME, LOG_PID | LOG_NDELAY | LOG_NOWAIT, LOG_DAEMON);
62     syslog(LOG_INFO, "staring sampled");
63
64     // We really don't want to take over syslogd or initd, so
65     // fork.
66     pid_t pid = fork();
67
68     // Well something went wrong. fork() uses standard unix
69     // errno functionality, so let's log the problem and exit.
70     if (pid < 0) {
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 58% 70:1

# Getting Started

```
64 // We really don't want to take over syslogd or initd, so
65 // fork.
66 pid_t pid = fork();
67
68 // Well something went wrong. fork() uses standard unix
69 // errno functionality, so let's log the problem and exit.
70 if (pid < 0) {
71     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
72     return ERR_FORK;
73 }
74
75 // We receive a PID if we're the parent process. If we're then
76 // parent, let's just exit. We only care about the forked child.
77 if (pid > 0) {
78     return OK;
79 }
80
81 // I'd like to be the leader of the session. If I can't,
82 // I'm out.
83 if (setsid() < -1) {
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 69% 83:1

# Forking

```
80
81 // I'd like to be the leader of the session. If I can't,
82 // I'm out.
83 if(setsid() < -1) {
84     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
85     return ERR_SETSID;
86 }
87
88 // We have console to write to anymore, so these file pointers
89 // are just silly. Close them.
90 close(STDIN_FILENO);
91 close(STDOUT_FILENO);
92 close(STDERR_FILENO);
93
94 // Are we creating files? Well, not in this example, but let's
95 // go ahead and set a sane UMASK anyway. This will give us
96 // read/write, and everybody else gets read permissions.
97 umask(S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
98
99 // What's our working directory going to be? Let's make it the root
```

NORMAL master > ./sample\_daemon.c unix < utf-8 < c 66% 80:1

# Sessions and Output

```
97  umask(S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
98
99  // What's our working directory going to be? Let's make it the root
100 // directory just to make navigation easier if we needed it.
101 if (chdir("/") < 0) {
102     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
103     return ERR_CHDIR;
104 }
105
106 // Now, we can only be controlled by signals as we have no interactive
107 // user session. So let's handle those signals, yeah?
108 signal(SIGTERM, _signal_handler);
109 signal(SIGHUP, _signal_handler);
110
111 // This is the daemon proces runloop. This could be a while statement,
112 // or some other construct, but daemons have specific responsibilities
113 // they need to handle, usually in a repeated way, and all that log
114 // would go here.
115 _do_work();
116
```

NORMAL master > ./sample\_daemon.c unix < utf-8 < c 89% 108:1

# Working Directories and Signals

```
25 } error_t;
26
27 // This is a signal handler. Signal is the signal passed to
28 // us to handle.
29 static void _signal_handler(const int signal) {
30     switch(signal) {
31         case SIGHUP:
32             break;
33         case SIGTERM:
34             syslog(LOG_INFO, "received SIGTERM, exiting.");
35             closelog();
36             exit(OK);
37             break;
38         default:
39             syslog(LOG_INFO, "received unhandled signal");
40     }
41 }
42
43 // This is where we do the work of the daemon. This example
44 // just counts and sleeps. Impressive!
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 21% 25:1

# Signal Handling



```
102     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
103     return ERR_CHDIR;
104 }
105
106 // Now, we can only be controlled by signals as we have no interactive
107 // user session. So let's handle those signals, yeah?
108 signal(SIGTERM, _signal_handler);
109 signal(SIGHUP, _signal_handler);
110
111 // This is the daemon proces runloop. This could be a while statement,
112 // or some other construct, but daemons have specific responsibilities
113 // they need to handle, usually in a repeated way, and all that log
114 // would go here.
115 _do_work();
116
117 // This should actually never be reached. We should be looping in
118 // _do_work() until the daemon is killed by a signal, at which
119 // point we exit the process.
120 return ERR_WTF;
121 }
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 100% 121:1

# Doing Actual Things!



```
38     default:
39         syslog(LOG_INFO, "received unhandled signal");
40     }
41 }
42
43 // This is where we do the work of the daemon. This example
44 // just counts and sleeps. Impressive!
45 static void _do_work(void) {
46     for(int i = 0; true; i++) {
47         syslog(LOG_INFO, "iteration: %d", i);
48         sleep(1);
49     }
50 }
51
52 int main(void) {
53     // Open syslog. We want to write directly to the log with
54     // no wait, and we want to include the PID of the daemon process.
55     // We are running a daemon, so set the LOG_DAEMON flag too.
56     //
57     // Why open here? well, if we have an error forking, we want to
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 47% 57:1

# Run Loop

# IoT: Client Devices

Configuring OS for Daemons

# Starting Daemons

## INIT, UPSTART, SYSTEMD

---

- ▶ Init was first, it's oldest - initially in Unix System V
- ▶ Upstart was next, used in Debian distros and Redhat
- ▶ Systemd is the most current service manager

## SYSTEMD CRITICISM

---

- ▶ Complex, violates Unix design philosophy
- ▶ But still widely used

Our IoT system uses init.

```
/etc
[# ls
fstab      hosts      mtab      passwd-   random-seed shadow-
group      init.d     network   profile   resolv.conf ssh
group-     inittab    os-release profile.d  services  ssl
hostname   issue     passwd    protocols shadow
[# ls init.d
S01logging S20urandom S40network S50sshd    rcK      rcS
#
```

# File Structure

init uses filenames to determine startup ordering

```
[
# Startup the system
[::sysinit:/bin/mount -t proc proc /proc
::sysinit:/bin/mount -o remount,rw /
::sysinit:/bin/mkdir -p /dev/pts
::sysinit:/bin/mkdir -p /dev/shm
::sysinit:/bin/mount -a
::sysinit:/bin/hostname -F /etc/hostname
# now run any rc scripts
::sysinit:/etc/init.d/rcS

# Put a getty on the serial port
ttyAMA0::respawn:/sbin/getty -L ttyAMA0 0 vt100 # GENERIC_SERIAL

# Stuff to do for the 3-finger salute
#::ctrlaltdel:/sbin/reboot

# Stuff to do before rebooting
::shutdown:/etc/init.d/rck
::shutdown:/sbin/swapoff -a
::shutdown:/bin/umount -a -r
```

# inittab

Things we do when the system starts up

```
[
# Start all init scripts in /etc/init.d
# executing them in numerical order.
#
for i in /etc/init.d/S??* ;do

    # Ignore dangling symlinks (if any).
    [ ! -f "$i" ] && continue

    case "$i" in
        *.sh)
            # Source shell script for speed.
            (
                trap - INT QUIT TSTP
                set start
                . $i
            )
            ;;
        *)
            # No sh extension, so fork subprocess.

```

# rcS

Starting all services on system start

```

[# Stop all init scripts in /etc/init.d
[# executing them in reversed numerical order.
[#
for i in $(ls -r /etc/init.d/S??*) ;do
[
[   # Ignore dangling symlinks (if any).
[   [ ! -f "$i" ] && continue

case "$i" in
    *.sh)
[       # Source shell script for speed.
[       (
[           trap - INT QUIT TSTP
[           set stop
[           . $i
[       )
[       ;;
[   *)
[       # No sh extension, so fork subprocess.

```

# rck

Stopping all services on shutdown

# IoT: Client Devices

An Example Startup Script



# Particular Script Format

#!/bin/sh

<stuff that always happens>

<start function>

<stop function>

<case handling args>

exit \$?

```
1 #!/bin/sh
2
3 DAEMON_NAME="DAEMON_NAME"
4
5 start() {
6     printf "Starting $DAEMON_NAME: "
7     echo "OK"
8 }
9
10 stop() {
11     printf "Stopping $DAEMON_NAME: "
12     echo "OK"
13 }
14
15 restart() {
16     stop
17     start
18 }
19
20 case "$1" in
21     start)
22         start
23         ;;
24     stop)
25         stop
26         ;;
27     restart|reload)
28         restart
29         ;;
30     *)
31         echo "Usage: $0 {start|stop|restart}"
32         exit 1
33 esac
34
35 exit $?
```

```
Usage: ./sample_framed {start|stop|restart}
[cclamb@ubuntu:~ $ ./sample_framed start
Starting DAEMON NAME: OK
[cclamb@ubuntu:~ $ ./sample_framed stop
Stopping DAEMON NAME: OK
[cclamb@ubuntu:~ $ ./sample_framed restart
Stopping DAEMON NAME: OK
Starting DAEMON NAME: OK
cclamb@ubuntu:~ $
```

# Running the Frame

# Filling Out the Frame

Let's use our sample daemon

```
1 #!/bin/sh
2
3 DAEMON_NAME="sampled"
4
5 start() {
6     printf "Starting $DAEMON_NAME: "
7     /usr/sbin/$DAEMON_NAME
8     touch /var/lock/$DAEMON_NAME
9     echo "OK"
10 }
11
12 stop() {
13     printf "Stopping $DAEMON_NAME: "
14     killall $DAEMON_NAME
15     rm -f /var/lock/$DAEMON_NAME
16     echo "OK"
17 }
18
19 restart() {
20     stop
21     start
22 }
23
24 case "$1" in
25     start)
26         start
27         ;;
28     stop)
29         stop
30         ;;
31     restart|reload)
32         restart
33         ;;
34     *)
35         echo "Usage: $0 {start|stop|restart}"
36         exit 1
37 esac
38
39 exit $?
```

# Move sampled to image

## MOVE SAMPLED TO DEVICE

---

- ▶ `scp -P 2222 sampled <username>@localhost:~/`

## MOVE IT TO /USR/SBIN

---

- ▶ ...this is the location pointed to in start script

## MOVE START SCRIPT TO DEVICE

---

- ▶ place in `/etc/init.d`
- ▶ I called `S80sampled`
- ▶ Reboot, login, and check out `/var/log/messages`

# IoT: Client Devices

Daemon Processes

# Daemon Processes

EVERY OS HAS SOMETHING LIKE THIS

---

- ▶ Called a service (windows-land) or a daemon (unix)
- ▶ I've seen both terms used with MacOS

STARTUP & SHUTDOWN

---

- ▶ Starts with the system, stops with the system

You'll use these in your project!

# Not a Regular Program

## OPERATING SYSTEM SERVICES

---

- ▶ Not all services are available (e.g. no console for output)

## FOLLOW CERTAIN CONVENTIONS

---

- ▶ Forking, default file permissions, group affiliation

## OPERATING SYSTEM CONFIGURATION

---

- ▶ Linux needs to be configured to start your daemon

# Languages

## MULTIPLE LANGUAGES SUPPORT DAEMONS

---

- ▶ Python and python-daemon
- ▶ Ruby with Daemons
- ▶ Perl, bash, C

## WE'RE EMBEDDED THOUGH

---

- ▶ No python, perl, ruby, etc.
- ▶ We do have bash and C



# Pros & Cons of C/Bash

## CONFIGURATION IN BASH

---

- (+) Well supported standard way to configure Linux startup
- (-) Really, the only supported way to configure Linux startup

## CODE DEVELOPMENT IN C

---

- (+) Native to Linux/Unix
- (+) Abundance of services and libraries
- (+) Well supported via build root
- (+) Small footprint
- (-) Wordy
- (-) Complex
- (-) Powerful

When you have a hammer...

# IoT: Client Devices

AXFS

# Background

## DEVELOPED BY JARED HULBERT

---

- Originally at Intel, now at Numonyx, released in 2008

## DESIGNED FOR XIP

---

- Execution-in-place, very handy for mobile/embedded
- Not integrated into Linux kernel well
- Clumsy patch for XIP in common use with CRAMFS

# Flash and XIP

## SO WHY XIP?


---

- ▶ Mobile and embedded systems are not desktop
- ▶ Flash memory is fast, disk is slow
- ▶ Running programs directly from Flash is possible
- ▶ Saves space, speeds up loading

## OBSTACLES

---

- ▶ Program loading is remarkably different
  - ▶ Shared libraries, code relocation, other dynamic loading




# Design - Mounting

Split Mounting

**NOR**: XIP Programs **NAND**: Media Storage

### **Superblock**

Volume information and offsets to region descriptors



### **Region Descriptor**

Contains an offset to a region, compression info, region size



### **Region**

ByteTable or data nodes (XIP, compressed, or byte aligned)

# Design - Data

(See Hulbert's *Introducing the Advanced XIP File System* for more details)

# IoT: Client Devices

ext2

# Background

## TYPICAL LINUX GENERAL PURPOSE FILESYSTEMS

---

- ▶ Remy Card, 1993
- ▶ Basic design in use today (see: ext4)
- ▶ Non-journaling
- ▶ Based on earlier Berkeley Fast Filesystem

## DRIVERS FOR EVERY MAJOR OS

---

- ▶ Though they may be third-party



# Aside: Journaling FS

## LIKE A DATABASE OF CHANGES

---

- ▶ Write to changes to *journal* before committing to disk
- ▶ *File writes* usually require many separate *disk writes*
- ▶ Non-atomicity leads to potential corruption

## CRASH RECOVERY

---

- ▶ Requires FS walk if no journal, not always recoverable
- ▶ Journal played back to ensure that all changes are applied

### **Superblock**

Describes filesystem (@ byte 1024, 1024 bytes in length)

---

### **Block Group Descriptor Table**

Immediately after superblock, locates block groups

1 ↓ 1..\*

### **Block Group**

Contains groups of blocks

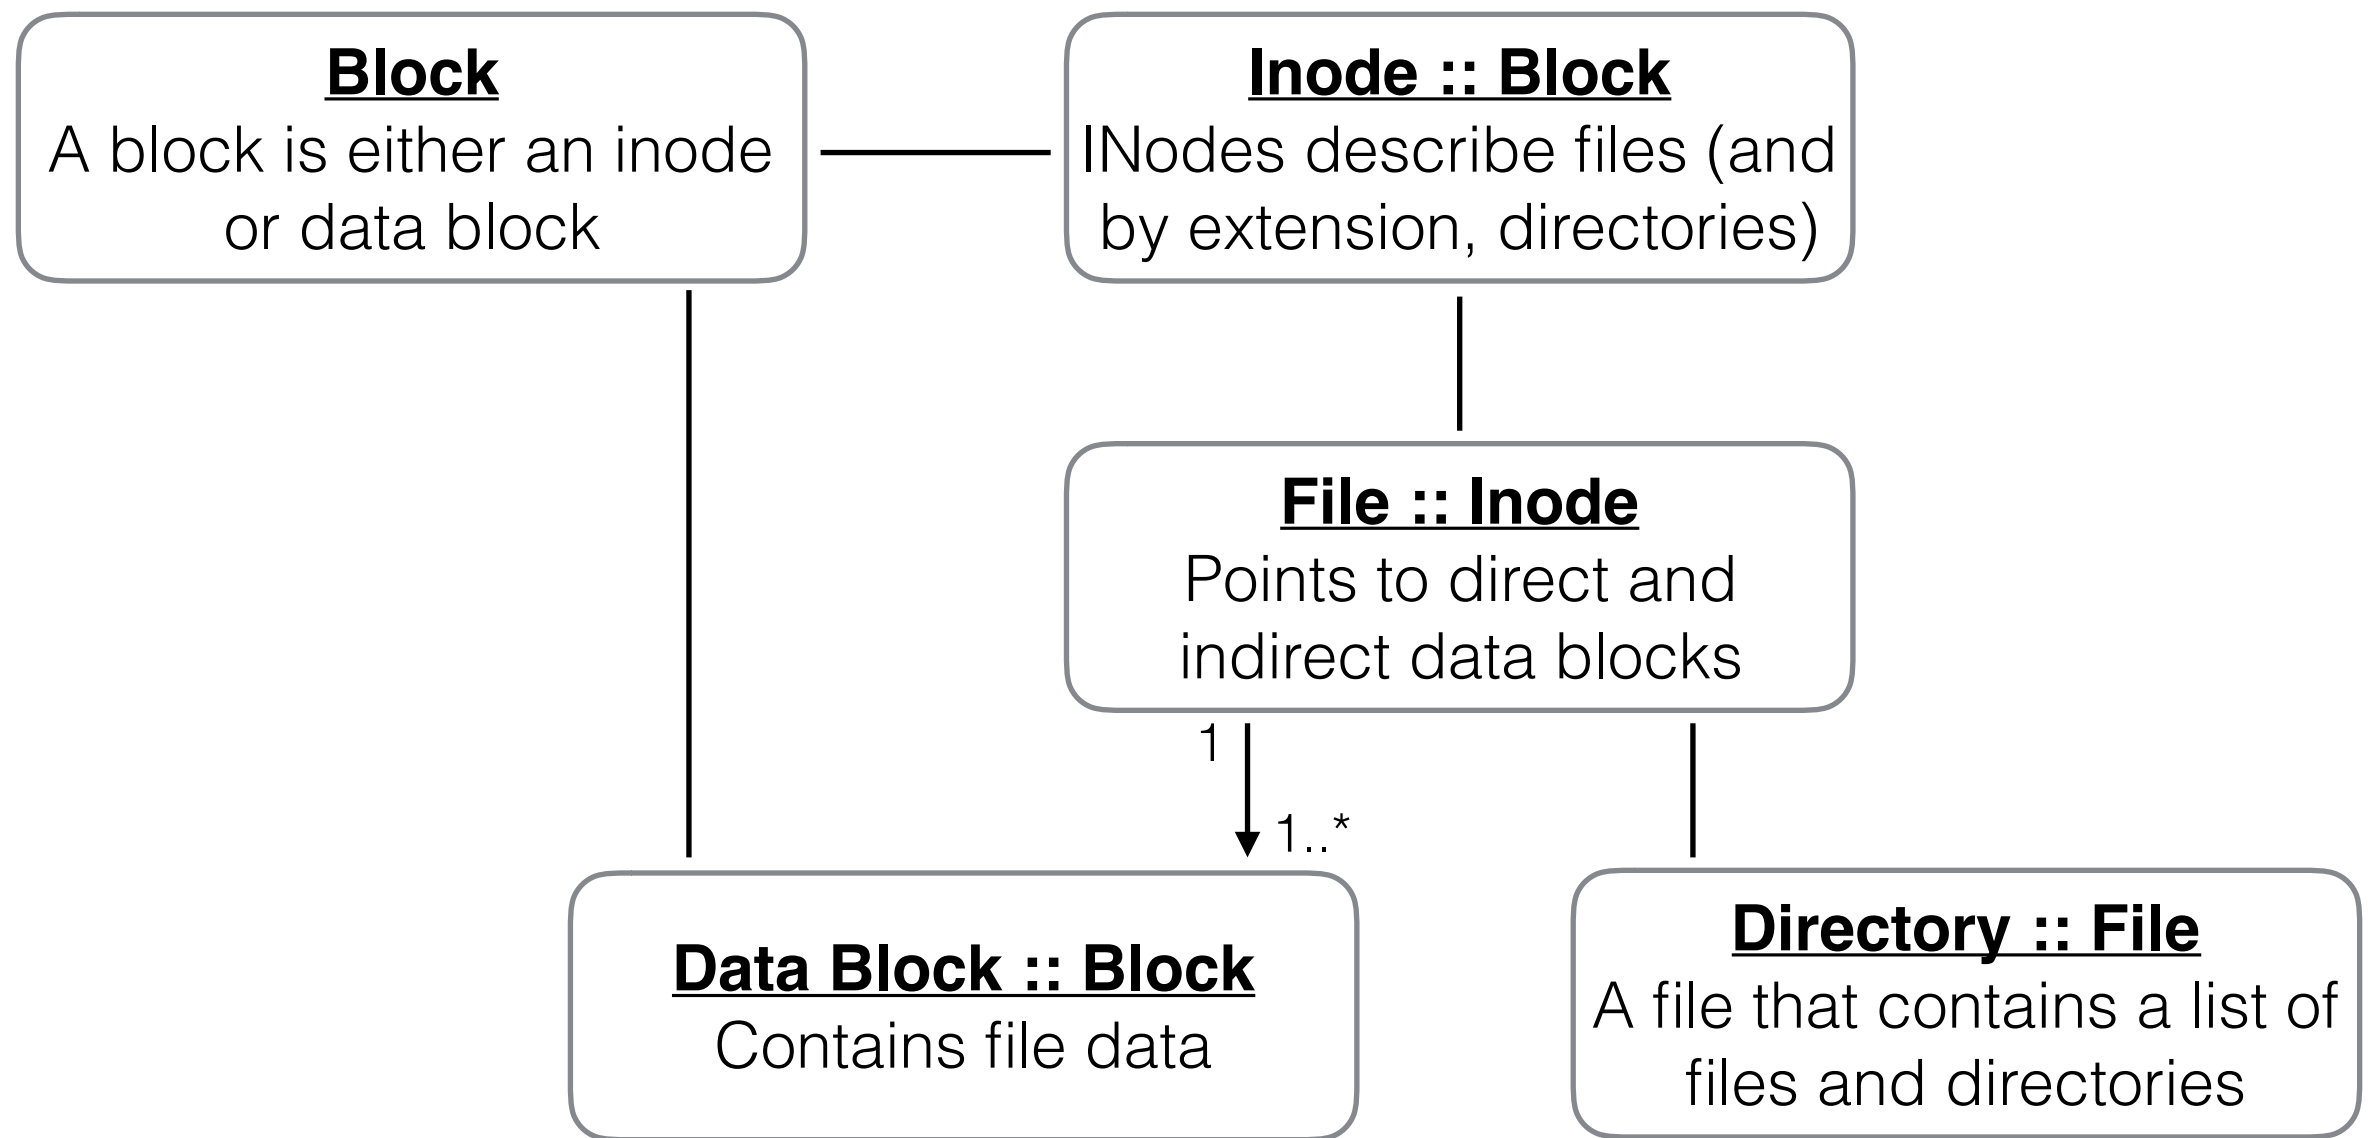
1 ↓ 1..\*

### **Block**

An inode or a data block

# Design - Blocks

Filesystem is based on blocks and block management  
***block size not tied to physical blocks***



# Design - Files and Directories

# IoT: Client Devices

Filesystems

## Filesystem images

```
[ ] axfs root filesystem
[ ] cloop root filesystem for the target device
[ ] cpio the root filesystem (for use as an initial RAM filesystem)
[ ] cramfs root filesystem
[*] ext2/3/4 root filesystem
    ext2/3/4 variant (ext2 (rev0)) --->
        () filesystem label
        (0) exact size in blocks (leave at 0 for auto calculation)
        (0) exact number of inodes (leave at 0 for auto calculation)
        (0) extra size in blocks
        (0) extra inodes
        (0) reserved blocks percentage
        Compression method (no compression) --->
[ ] initial RAM filesystem linked into linux kernel
[ ] jffs2 root filesystem
[ ] romfs root filesystem
[ ] squashfs root filesystem
[ ] tar the root filesystem
[ ] ubifs root filesystem
[ ] yaffs2 root filesystem
```

# Lots of Options

so what do you choose? We'll take it from the top.

**Always pay attention to licensing**

# AXFS

## WHAT IS IT?

---

- ▶ Advanced XIP Filesystem
- ▶ Compressed, read-only, execute-in-place

## WHY WOULD YOU USE IT?

---

- ▶ Fast boot and load, small footprint

## WHY WOULDN'T YOU USE IT?

---

- ▶ You need writeable filesystem
- ▶ Hardware support

# CLOOP

## WHAT IS IT?

---

- Compressed block driver, similar to Apple DMG
- Read-only block devices, transparent compression

## WHY WOULD YOU USE IT?

---

- Frequently used with Live CD
- Not a filesystem, but block device support (under FS)

## WHY WOULDN'T YOU USE IT?

---

- You want to use a full filesystem solution

# CRAMFS

## WHAT IS IT?

---

- ▶ Compressed ROM Filesystem, Read-only, Linux
- ▶ Simple, fast, small

## WHY WOULD YOU USE IT?

---

- ▶ Only if you're stuck with it; It's obsoleted by SquashFS

## WHY WOULDN'T YOU USE IT?

---

- ▶ Don't use it for new systems (unless forced for some reason)



# ext2/3/4

## WHAT IS IT?

---

- ▶ Very common Linux filesystem(s)
- ▶ ext2 non-journaling, used with Flash and SD cards

## WHY WOULD YOU USE IT?

---

- ▶ Writeable filesystem (moving programs to image, caching, etc.)
- ▶ Limit writes to storage (no journal to maintain)

## WHY WOULDN'T YOU USE IT?

---

- ▶ Writeable is more exploitable, speed

# RAM filesystem

## WHAT IS IT?

---

- ▶ A filesystem configured in RAM (i.e. a RAMDisk)

## WHY WOULD YOU USE IT?

---

- ▶ Very fast, gives initial filesystem while other loads

## WHY WOULDN'T YOU USE IT?

---

- ▶ Need more space than RAMDisk will provide
- ▶ Limit RAM usage

# JFFS2

## WHAT IS IT?

---

- ▶ Journaling Flash Filesystem

## WHY WOULD YOU USE IT?

---

- ▶ You want journaling, but you're using Flash memory
- ▶ You don't care about write degradation
- ▶ Compression

## WHY WOULDN'T YOU USE IT?

---

- ▶ There's successor filesystems (e.g. YAFFS)
- ▶ Slow boot, difficult filesystem analysis

# ROMFS

## WHAT IS IT?

---

- ▶ A very small, simple filesystem for EEPROMs

## WHY WOULD YOU USE IT?

---

- ▶ Kernel module storage

## WHY WOULDN'T YOU USE IT?

---

- ▶ If you don't know you need it, don't use it

# SquashFS

## WHAT IS IT?

---

- Successor to cramfs
- Compressed, read-only, large block support, low-overhead

## WHY WOULD YOU USE IT?

---

- You want a modern, read-only, compressed filesystem
- You don't care about XIP

## WHY WOULDN'T YOU USE IT?

---

- You want a writeable filesystem
- You care about XIP

# UBIFS

## WHAT IS IT?

---

- ▶ Unsorted block image filesystem
- ▶ Successor to JFFS2

## WHY WOULD YOU USE IT?

---

- ▶ Better than JFFS2 for large NAND Flash
- ▶ Better failure tolerance, compression support

## WHY WOULDN'T YOU USE IT?

---

- ▶ Locked into JFFS2, Hardware limitations

# YAFFS2

## WHAT IS IT?

---

- ▶ Yet Another Flash Filesystem
- ▶ Log structured, high data-integrity goals

## WHY WOULD YOU USE IT?

---

- ▶ Portable, Fast, supports modern hardware

## WHY WOULDN'T YOU USE IT?

---

- ▶ Hardware restrictions or licensing
- ▶ No compression

# IoT: Client Devices

Embedded Kernels



# Embedded Kernels

I'm not going to discuss the Linux kernel in depth.

# Guidelines for Kernel

## YOU MAY NOT HAVE MUCH CHOICE

---

- ▶ The kernel modules you need are dictated by hardware
- ▶ You'll likely have a fixed config file

## BUILDROOT GIVES YOU SOME OPTIONS

---

- ▶ For the most part, use the defaults

## Kernel

### [\*] Linux Kernel

```
    Kernel version (Custom version) --->
    (4.8.1) Kernel version
    () Custom kernel patches
    Kernel configuration (Using a custom (def)config file) --->
    (board/qemu/arm-versatile/linux-4.8.config) Configuration file path
    () Additional configuration fragment files
    Kernel binary format (zImage) --->
    Kernel compression format (gzip compression) --->
[*] Build a Device Tree Blob (DTB)
    Device tree source (Use a device tree present in the kernel.)
    (versatile-pb) Device Tree Source file names
[ ] Install kernel image to /boot in target
    Linux Kernel Extensions --->
    Linux Kernel Tools --->
```

# Changing Defaults

If you do change defaults...  
**...back up your config(s) first!**

```
[cclamb@ubuntu:~/buildroot-2016.11.1 $ head -20 board/qemu/arm-versatile/linux-4.8.config
CONFIG_SYSVIPC=y
CONFIG_MODULES=y
CONFIG_MODULE_UNLOAD=y
# CONFIG_ARCH_MULTI_V7 is not set
CONFIG_ARCH_VERSATILE=y
CONFIG_PCI=y
CONFIG_PCI_VERSATILE=y
CONFIG_AEABI=y
CONFIG_NET=y
CONFIG_PACKET=y
CONFIG_UNIX=y
CONFIG_INET=y
CONFIG_SCSI=y
CONFIG_BLK_DEV_SD=y
CONFIG_SCSI_SYM53C8XX_2=y
CONFIG_NETDEVICES=y
CONFIG_8139CP=y
CONFIG_PHYLIB=y
CONFIG_INPUT_EVDEV=y
CONFIG_SERIO_AMBAKMI=y
cclamb@ubuntu:~/buildroot-2016.11.1 $
```

# Defconfig Files

# IoT: Client Devices

Bootloaders

# We're Not Using One!

## WE'RE RUNNING A SIMULATOR

---

- ▶ So we're not using one
- ▶ In real life you'll need one

## BOOT LOADERS

---

- ▶ DAS U-Boot, gummiboot, AT91 Bootstrap, etc
- ▶ Take a look at the Bootloader menu

# Bootloaders

Yeah, these are complex too.

# Bootloader does What?

## LOADS THE OPERATING SYSTEM

---

- ▶ Transitions control to OS

## HANDLES INITIAL CONFIGURATION

---

- ▶ Power on self tests first...
- ▶ ...then loads bootloader...
- ▶ ...then loads OS.

## COMMONLY USE MULTIPLE STAGES

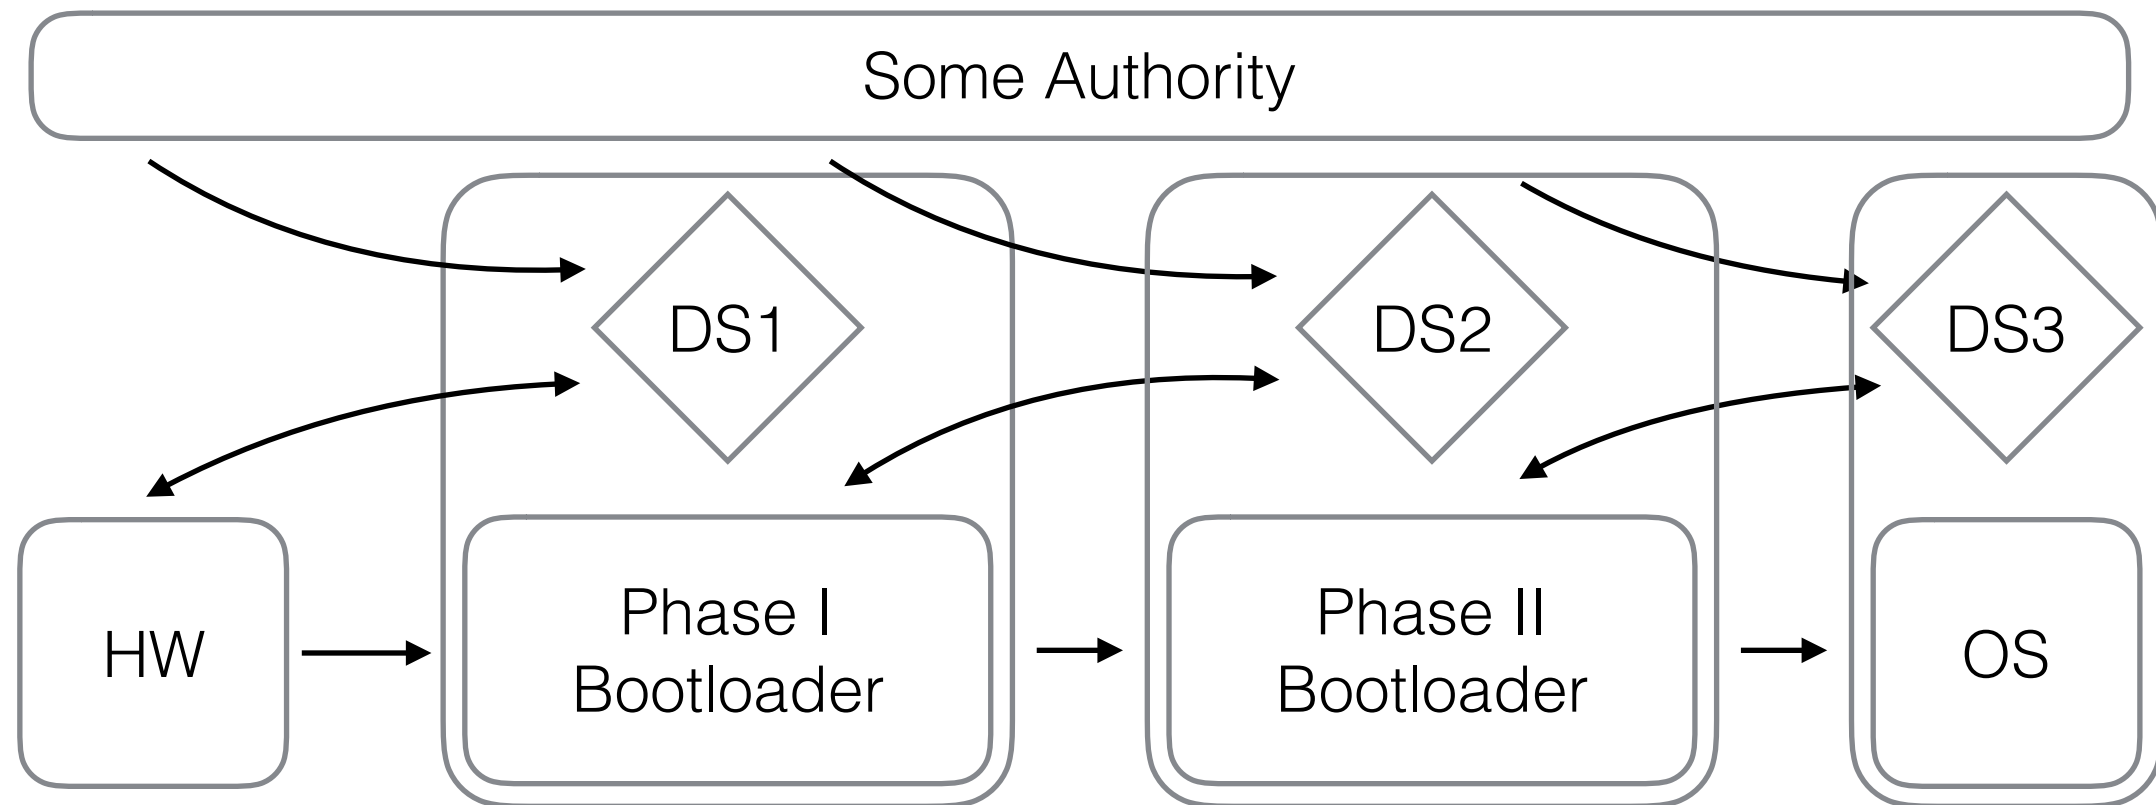
---

- ▶ Different loaders of different complexity



# IoT: Client Devices

Bootloaders



# Secure Boot

How do we do this?

# IoT: Client Devices

Securing the Kernel and OS

# Securing the Kernel

OKAY, YOU'RE NOT A KERNEL/OS DEVELOPER

---

- ▶ You're not going too be writing secure kernel code

YOU ARE RESPONSIBLE FOR A SECURE KERNEL

---

- ▶ ...on your device

So what does this mean?

# Use a Secure Kernel

## DON'T USE BUGGY CODE

---

- ▶ This includes kernels, libraries, etc.

## OLD? KNOWN ISSUES?

---

- ▶ Don't use it
- ▶ Use the newest most secure code you can

# Ship Secure Tooling

## OS IMAGES COME WITH LOTS OF STUFF

---

- Only ship what you need
- Don't ship things you don't

## ONGOING MAINTENANCE AND ANALYSIS

---

- This is a real need too
- Make sure anything that you ship on your device is as secure as possible!

## DON'T WRITE YOUR OWN

---

- Don't create your own protocols or encryption

# Build Updatable Devices

VULNERABILITIES **WILL** POP UP

---

- ▶ They always do.

MAKE SURE THAT WHEN THEY DO YOU CAN FIX THEM

---

- ▶ You'll need to do this at scale
- ▶ You'll need to secure this too

# IoT: Client Devices

Securing your Application



# Yes, you.

## SECURING KERNELS, OS

---

- ▶ You're responsible for this
- ▶ You don't have much control

You have total control over **YOUR** software.

# Know Typical Flaws

## KNOW TYPICAL SOFTWARE FLAWS

---

- ▶ C: buffer overflows, integer underflows, bad coding, etc.
- ▶ Web: encryption, authentication, authorization, etc.

## KNOW RESOURCES

---

- ▶ C secure programming standard
- ▶ SANS top 20
- ▶ OWASP secure programming practices

# Red Teams & Pentest

## CODE REVIEWS

---

- Review code with development colleagues

## RED TEAMS

---

- Have adversarial design and code reviews

## ATTACK!

---

- Penetration test the final product

# Don't DYI

## OTHER EYES == BETTER

---

- ▶ Find others to review your code
- ▶ Other organizations to red team and pentest

## STUDY SECURE CODING

---

- ▶ It's not hard, you just need to do it!