

# IoT: Client Devices

A Sample Daemon

```
51
52 int main(void) {
53     // Open syslog. We want to write directly to the log with
54     // no wait, and we want to include the PID of the daemon process.
55     // We are running a daemon, so set the LOG_DAEMON flag too.
56     //
57     // Why open here? well, if we have an error forking, we want to
58     // register the error to syslog. To do that, we need an open log.
59     // The child process will inherit this log, as we don't close it
60     // in the parent when it exits, this opening here is A+.
61     openlog(DAEMON_NAME, LOG_PID | LOG_NDELAY | LOG_NOWAIT, LOG_DAEMON);
62     syslog(LOG_INFO, "staring sampled");
63
64     // We really don't want to take over syslogd or initd, so
65     // fork.
66     pid_t pid = fork();
67
68     // Well something went wrong. fork() uses standard unix
69     // errno functionality, so let's log the problem and exit.
70     if (pid < 0) {
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 58% 70:1

# Getting Started

```
64 // We really don't want to take over syslogd or initd, so
65 // fork.
66 pid_t pid = fork();
67
68 // Well something went wrong. fork() uses standard unix
69 // errno functionality, so let's log the problem and exit.
70 if (pid < 0) {
71     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
72     return ERR_FORK;
73 }
74
75 // We receive a PID if we're the parent process. If we're then
76 // parent, let's just exit. We only care about the forked child.
77 if (pid > 0) {
78     return OK;
79 }
80
81 // I'd like to be the leader of the session. If I can't,
82 // I'm out.
83 if (setsid() < -1) {
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 69% 83:1

# Forking

```
80
81 // I'd like to be the leader of the session. If I can't,
82 // I'm out.
83 if(setuid(0) < 0) {
84     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
85     return ERR_SETUID;
86 }
87
88 // We have console to write to anymore, so these file pointers
89 // are just silly. Close them.
90 close(STDIN_FILENO);
91 close(STDOUT_FILENO);
92 close(STDERR_FILENO);
93
94 // Are we creating files? Well, not in this example, but let's
95 // go ahead and set a sane UMASK anyway. This will give us
96 // read/write, and everybody else gets read permissions.
97 umask(S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
98
99 // What's our working directory going to be? Let's make it the root
```

NORMAL master > ./sample\_daemon.c unix < utf-8 < c 66% 80:1

# Sessions and Output

```
97  umask(S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
98
99  // What's our working directory going to be? Let's make it the root
100 // directory just to make navigation easier if we needed it.
101 if (chdir("/") < 0) {
102     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
103     return ERR_CHDIR;
104 }
105
106 // Now, we can only be controlled by signals as we have no interactive
107 // user session. So let's handle those signals, yeah?
108 signal(SIGTERM, _signal_handler);
109 signal(SIGHUP, _signal_handler);
110
111 // This is the daemon proces runloop. This could be a while statement,
112 // or some other construct, but daemons have specific responsibilities
113 // they need to handle, usually in a repeated way, and all that log
114 // would go here.
115 _do_work();
116
```

NORMAL master > ./sample\_daemon.c unix < utf-8 < c 89% 108:1

# Working Directories and Signals

```
25 } error_t;
26
27 // This is a signal handler. Signal is the signal passed to
28 // us to handle.
29 static void _signal_handler(const int signal) {
30     switch(signal) {
31         case SIGHUP:
32             break;
33         case SIGTERM:
34             syslog(LOG_INFO, "received SIGTERM, exiting.");
35             closelog();
36             exit(OK);
37             break;
38         default:
39             syslog(LOG_INFO, "received unhandled signal");
40     }
41 }
42
43 // This is where we do the work of the daemon. This example
44 // just counts and sleeps. Impressive!
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 21% 25:1

# Signal Handling



```
102     syslog(LOG_ERR, ERROR_FORMAT, strerror(errno));
103     return ERR_CHDIR;
104 }
105
106 // Now, we can only be controlled by signals as we have no interactive
107 // user session. So let's handle those signals, yeah?
108 signal(SIGTERM, _signal_handler);
109 signal(SIGHUP, _signal_handler);
110
111 // This is the daemon proces runloop. This could be a while statement,
112 // or some other construct, but daemons have specific responsibilities
113 // they need to handle, usually in a repeated way, and all that log
114 // would go here.
115 _do_work();
116
117 // This should actually never be reached. We should be looping in
118 // _do_work() until the daemon is killed by a signal, at which
119 // point we exit the process.
120 return ERR_WTF;
121 }
```

NORMAL master > ./sample daemon.c unix < utf-8 < c 100% 121:1

# Doing Actual Things!

```

38     default:
39         syslog(LOG_INFO, "received unhandled signal");
40     }
41 }
42
43 // This is where we do the work of the daemon. This example
44 // just counts and sleeps. Impressive!
45 static void _do_work(void) {
46     for(int i = 0; true; i++) {
47         syslog(LOG_INFO, "iteration: %d", i);
48         sleep(1);
49     }
50 }
51
52 int main(void) {
53     // Open syslog. We want to write directly to the log with
54     // no wait, and we want to include the PID of the daemon process.
55     // We are running a daemon, so set the LOG_DAEMON flag too.
56     //
57     // Why open here? well, if we have an error forking, we want to

```

NORMAL master > ./sample daemon.c unix < utf-8 < c 47% 57:1

# Run Loop